

SATYA AKHIL GALLA

Boston, MA

+1-857-540-4550 | akhilg@bu.edu | [LinkedIn](#) | [GitHub](#) | [Portfolio](#) | [Scholar](#)

EDUCATION

Boston University
Masters in Artificial Intelligence

Sep 2024 – Dec 2025
GPA: 3.7/4.0

IIIT Sri City
B.Tech in Computer Science

Dec 2020 – May 2024
GPA: 8.0/10

EXPERIENCE

Aaizel International Tech Pvt Ltd
Deep Learning Intern

Feb 2024 – May 2024
India

- Built the company's first scene-graph pipeline for IR + RGB satellite imagery: adapted ReTR to identify objects and relations.
- Deployed a Gradio web app that lets analysts upload a satellite tile and instantly view its semantic graph, accelerating remote-sensing intelligence workflows.

Terrafic Inc
Machine Learning Research Intern

June 2023 – Nov 2023
India

- Built the startup's first satellite-imagery ML pipeline; implemented 4× / 6× super-resolution for a defense-mapping demo.
- Added SAM-based segmentation of maritime and urban features, delivering a live geospatial proof-of-concept for founders.

ISRO - Indian Space Research Organisation
Machine Learning Research Intern

Dec 2022 – May 2023
India

- Built a Graph-Convolution pipeline for Chandrayaan-1/2 hyperspectral data and research on adaptive graph strategies.

Spatial Analytics and Machine Intelligence(SAMI) Lab - IIIT Sri City | [Lab Website](#)
Undergraduate Researcher

Aug 2022 – May 2024
India

- Developed graph-based models that raised lunar hyperspectral-image classification to 98 % accuracy.
- Authored an IEEE WHISPERS 2023 paper and NASA NESF 2023 poster; work funded by ISRO.

PROJECTS

Text-to-Music Generation with Transformers and Diffusion Models | *PyTorch, HuggingFace* | [GitHub](#)

Jan 2025 – May 2025

- First comparative study of transformer-first vs. diffusion-first cascades; showed MusicGen boosts semantic alignment while AudioLDM enriches audio detail—complementary, non-overlapping strengths.
- On the MusicCaps dataset, the best AudioLDM→MusicGen cascade kept semantic match (CLAP **0.35**) and cut audio distance (FAD 6.68 → **6.55**), outperforming both standalone models.

Layer-Residual Co-Attention Networks (LRCN) for VQA | *PyTorch* | [GitHub](#)

Jan 2025 – May 2025

- Open-sourced the first full PyTorch replication of all three LRCN variants and added a cross-attention visualisation tool.
- Matched the paper's VQA-v2 result and extending to CLEVR in upcoming work

Multimodal Skin Cancer Detection | *PyTorch, TensorFlow* | [GitHub](#)

Sep 2024 – Dec 2024

- Handled extreme class imbalance (0.1%) with resampling and weighted loss; trained image and tabular models separately.
- Fused both views in a voting ensemble raising ROC-AUC to **0.96**— 15 pts higher than the best single-modal baseline.

Minimizing Hallucination in Medical LLMs | *LLaMa, LangChain, FAISS* | [GitHub](#)

Sep 2024 – Dec 2024

- Built a medical RAG stack indexing 207 K verified Q-A pairs to ground responses and curb hallucinations.
- Raised SBERT sentence similarity from 0.63 to **0.87** (+38 %) over vanilla LLaMA-3, delivering more accurate patient-care answers.

PUBLICATIONS

Efficient Graph Formulation and Latent Space Integration for Lunar Hyperspectral Image Classification | *IEEE WHISPERS 2023, Athens* | [Publication](#)

- First study on Chandrayaan-1 and 2 hyperspectral data: benchmarked six graph-construction schemes; baseline handcrafted cosine graph scored 94 % accuracy, while a non-linear spectral-spatial graph lifted GCN accuracy to **97 %** (+3 pp).
- Integrated autoencoder latents with the best graph, pushing accuracy to **98.4 %** (+1.4 pp) and forming the core of an ISRO-funded mineral-mapping pipeline.

Enhancing Hyperspectral Classification through GCNs with Adaptive Graph Construction | *NASA NESF 2023 Poster #38, Maryland* | [Poster](#)

- First work on Chandrayaan-2 IIRS data to compare six graph-construction schemes for GCNs; introduced a non-linear spectral-spatial graph that lifted overall accuracy from 91 % to **97 %**.
- Showed adaptive graphs outperform handcrafted layouts and provide more interpretable mineral maps—insight.

SKILLS

Frameworks: PyTorch, TensorFlow, Keras, Scikit-learn, LangChain

Languages: Python, C++, SQL, JavaScript

Tools & Libraries: Git, FAISS, HuggingFace, OpenCV, Bash

Data & Visualization: NumPy, Pandas, Matplotlib, Seaborn